U.S. Traffic Accidents

Data Science Career Track

Sara Satti

August 2020

Objective:

Predict road accident severity given weather conditions, location, points of interest, and other engineered features.

Outline

- Data Wrangling
  - Dataset
  - Data Manipulation
    - Missing Data Handling
    - Feature Engineering
- Exploratory Data Analysis
- Modeling
- Conclusions and Recommendations

Data Wrangling

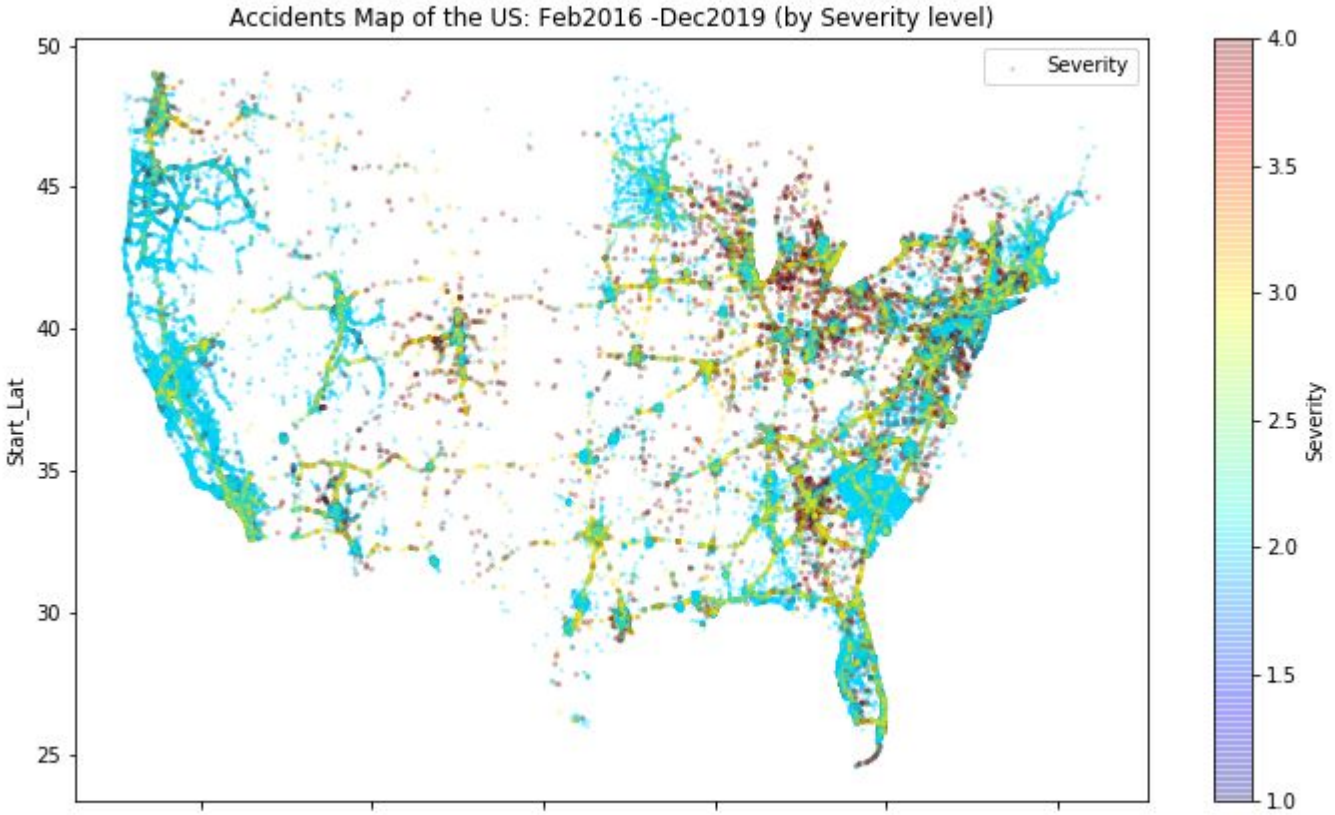

Dataset:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400000 entries, 0 to 399999
Data columns (total 49 columns):
ID                      400000 non-null object
Source                  400000 non-null object
TMC                     301701 non-null float64
Severity                400000 non-null int64
Start_Time              400000 non-null object
End_Time                400000 non-null object
Start_Lat               400000 non-null float64
Start_Lng               400000 non-null float64
End_Lat                 98299 non-null float64
End_Lng                 98299 non-null float64
Distance(mi)            400000 non-null float64
Description             400000 non-null object
Number                  141503 non-null float64
Street                  400000 non-null object
Side                    400000 non-null object
City                    399990 non-null object
County                  400000 non-null object
State                   400000 non-null object
Zipcode                 399870 non-null object
Country                 400000 non-null object
Timezone                399552 non-null object
Airport_Code            399213 non-null object
Weather_Timestamp       394983 non-null object
Temperature(F)          392443 non-null float64
Wind_Chill(F)           151072 non-null float64
Humidity(%)             392044 non-null float64
Pressure(in)            393526 non-null float64
Visibility(mi)          391156 non-null float64
Wind_Direction          393897 non-null object
Wind_Speed(mph)         341003 non-null float64
Precipitation(in)       131252 non-null float64
Weather_Condition       391170 non-null object
Amenity                 400000 non-null bool
Bump                    400000 non-null bool
Crossing                400000 non-null bool
Give_Way                400000 non-null bool
Junction                400000 non-null bool
No_Exit                 400000 non-null bool
Railway                 400000 non-null bool
Roundabout              400000 non-null bool
Station                 400000 non-null bool
Stop                    400000 non-null bool
Traffic_Calming         400000 non-null bool
Traffic_Signal          400000 non-null bool
Turning_Loop            400000 non-null bool
Sunrise_Sunset          399989 non-null object
Civil_Twilight          399989 non-null object
Nautical_Twilight       399989 non-null object
Astronomical_Twilight   399989 non-null object
dtypes: bool(13), float64(14), int64(1), object(21)
memory usage: 114.8+ MB
```

Data Manipulation - 1

Missing data handling:

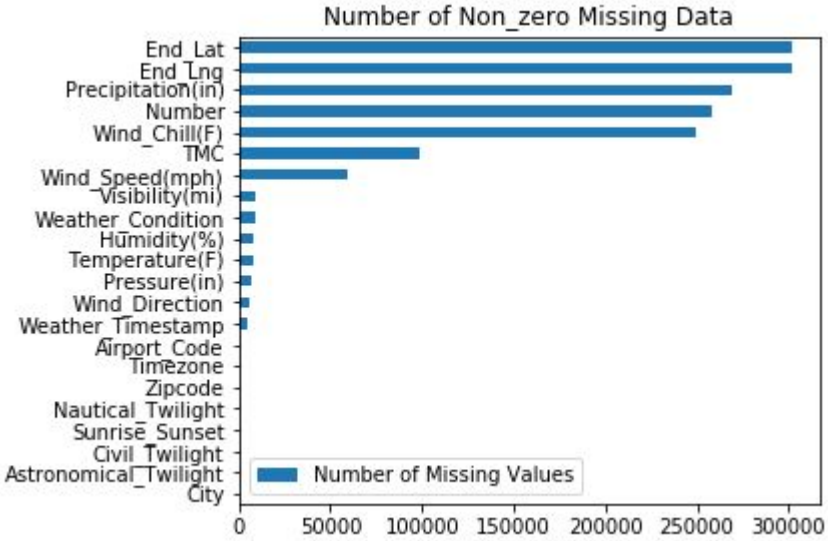

- ➔ isnull() count, and wasnull column addition.
- ➔ Missing precipitation = 0
- ➔ Start Lat/lng instead of end Lat/Lng.
- ➔ Wind_Chill: through linear regression.
- ➔ Others by infilling with Mean/median.

Data Manipulation - 2

Feature Engineering

1. Wasnull columns → tag null values within DF

from datetime data

2. Day_of_Week

3. Hour_of_Day

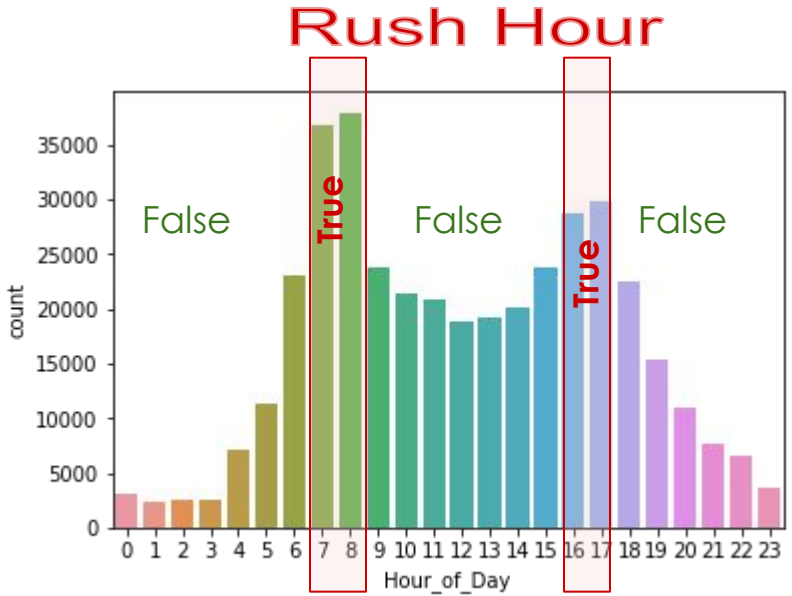

EDA

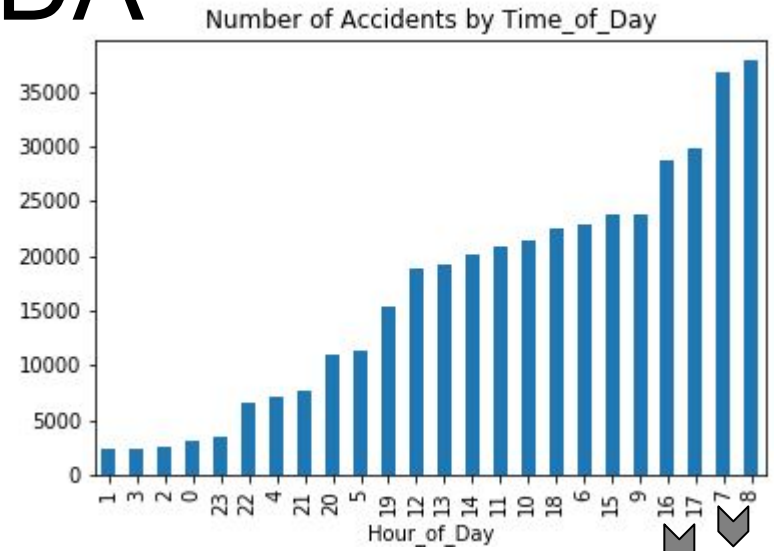

Morning R.H.

Evening Rush hour

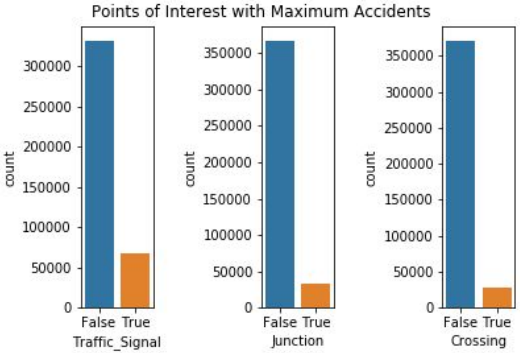

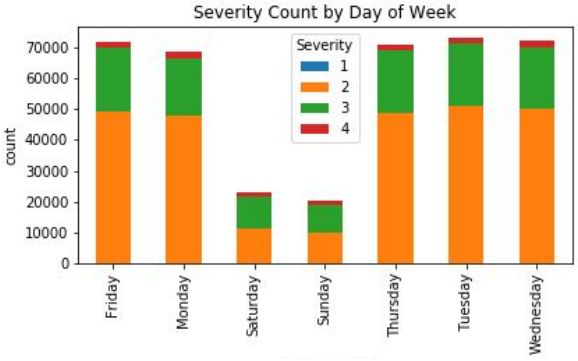

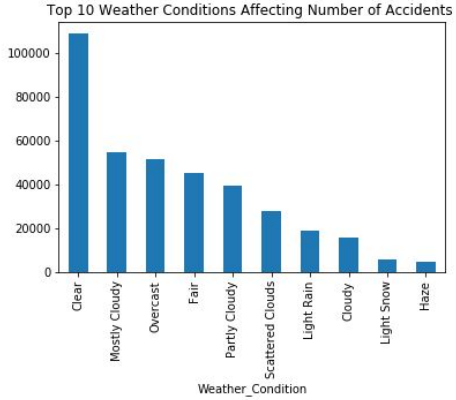

EDA

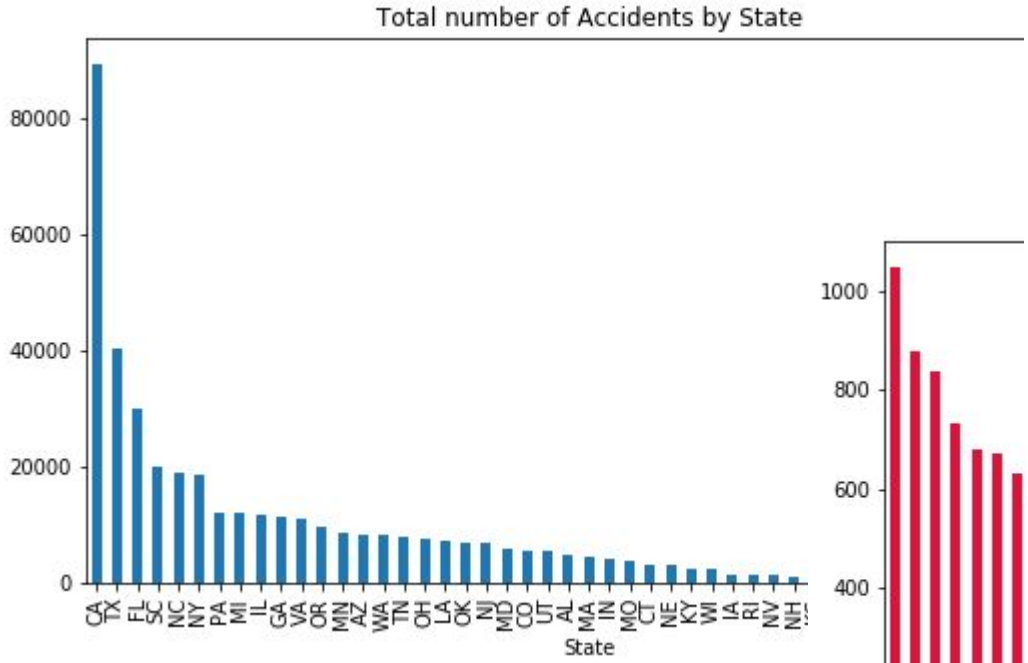

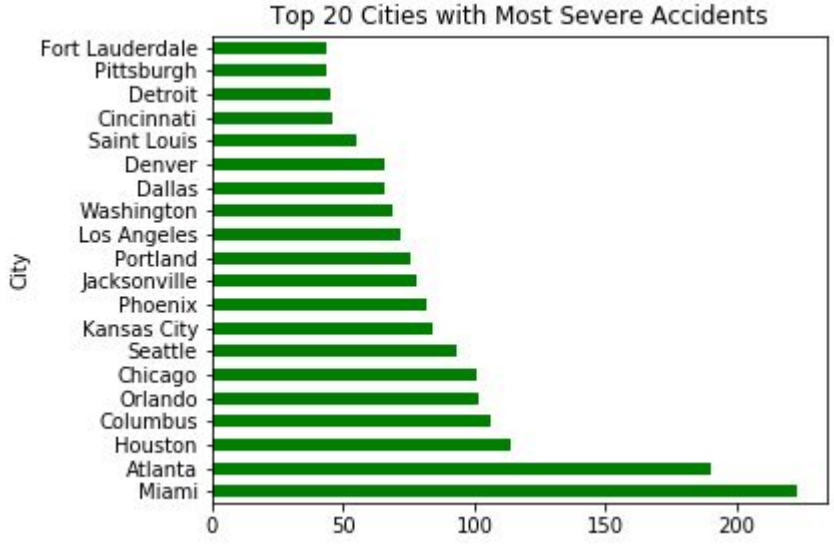

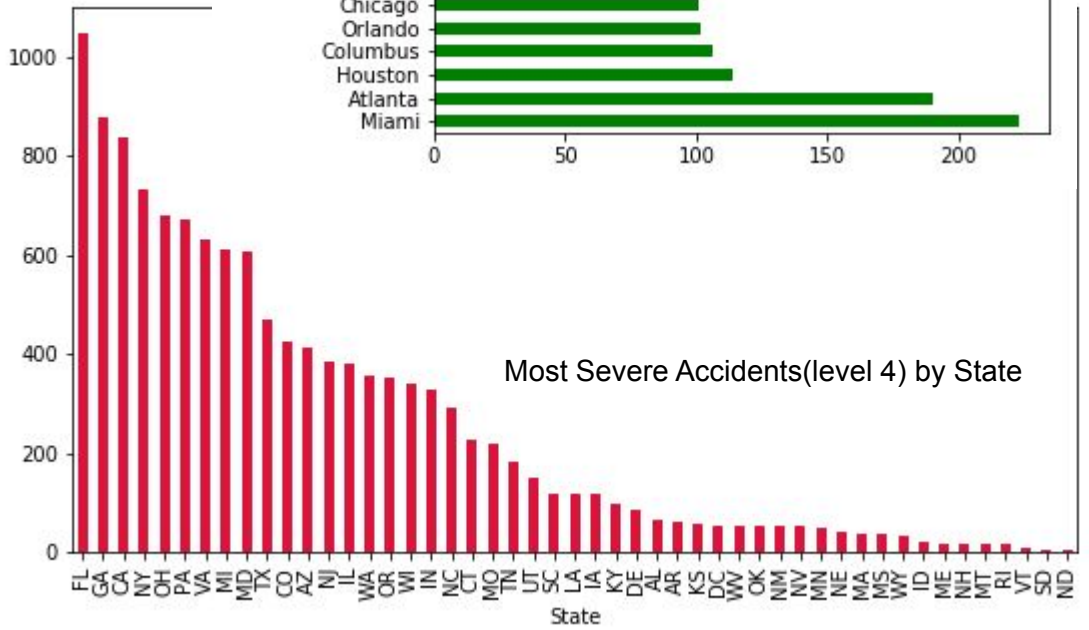

Most Severe Accidents(level 4) by State

Data Manipulation - 3

Modeling prep

subset 1
Lat/Long+Distance+weather

subset 2
subset 1 +
Traffic signal/junction/Rush_Hour/Side/Day_of_Week

subset 3
subset 2 + State

subset 4
subset 3 + All points of interest + wasnull

subset 5
subset 4 + drop wasnull

Re-sample (50k)

Modeling

SDG, KNN, Logistic Regression, Random Forest

Cross Validation No Scaling

Cross Validation **Scaling**

LogReg

Random Forest

Binary

MultiClass

Evaluation

Modeling

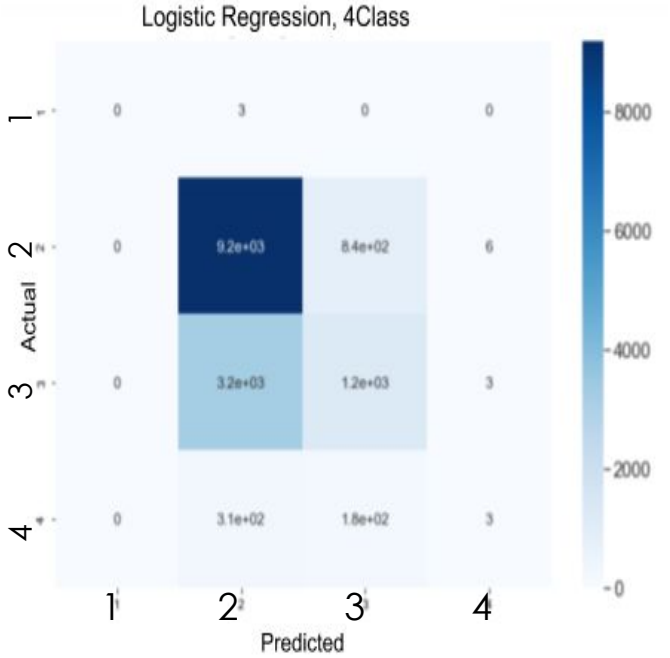

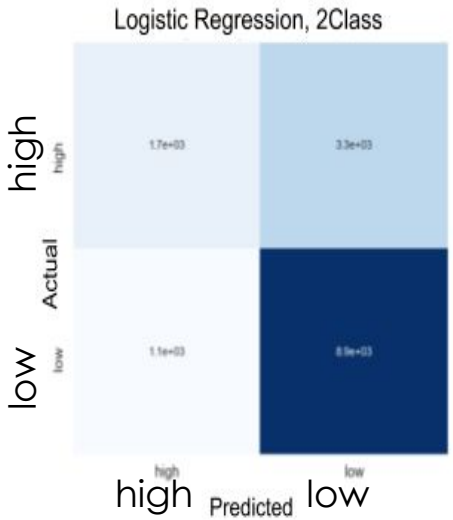

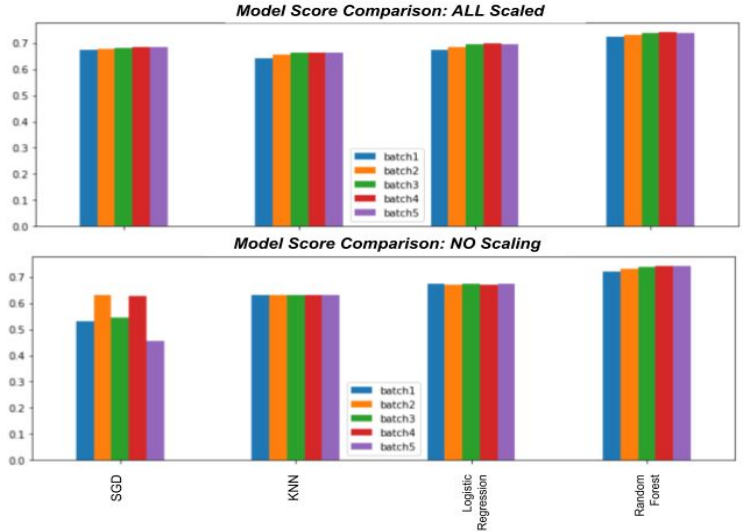

MultiClass and Binary Classification Reports

Logistic Regression

	precision	recall	f1-score	support
1	0	0	0	3
2	0.72	0.92	0.81	10024
3	0.55	0.28	0.37	4481
4	0.25	0.01	0.01	492
	precision	**recall**	**f1-score**	**support**
high	0.61	0.34	0.43	4973
low	0.73	0.89	0.8	10027

Modeling

Random Forest

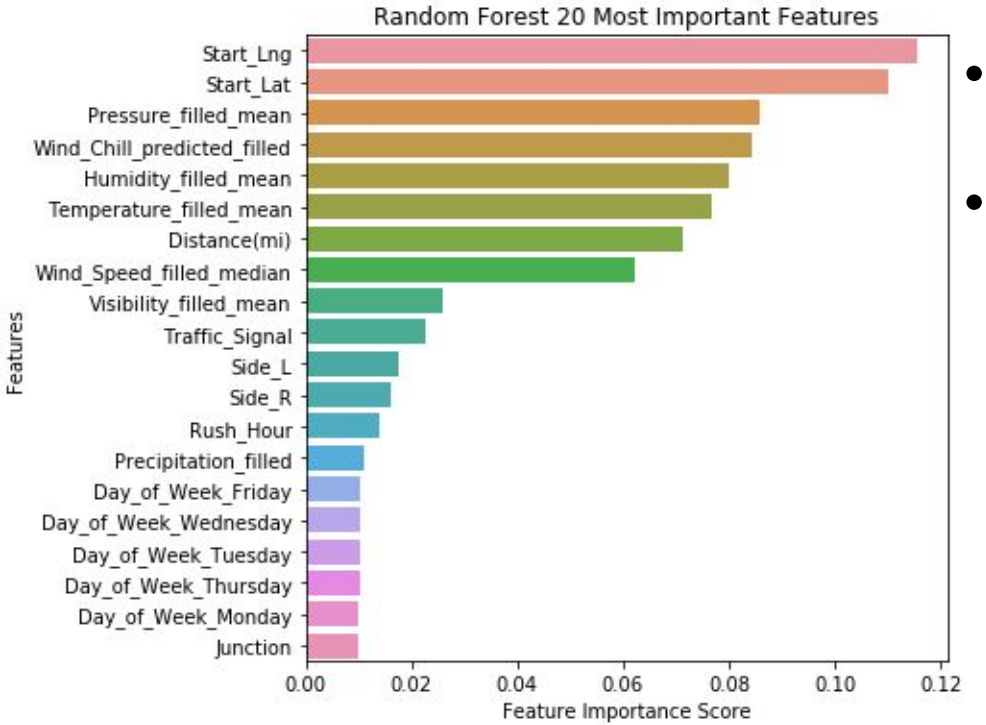

- **78% accuracy.**
- **f1 scores:**
  - **‘low’ = 85%**
  - **‘high’=78%**

Conclusions/ Recommendations

- Of the four machine learning algorithms used Random Forest gave the best results (78% accuracy).

- Room for improvements through:

more feature engineering and advanced machine learning.

The End

Q&A